

Explainable Confidence Algorithms for EHR QA LLM Outputs

1. Introduction

For EHR-based question answering, the most explainable algorithms for assigning a confidence score to an LLM's output are not raw token probabilities or simple self-reported confidence alone. The literature points instead to five practical families: consistency-based methods that compare multiple sampled answers, evidence-grounded methods that score support from retrieved EHR or biomedical evidence, verification-based methods that check the model's own reasoning or claims, post-hoc calibration methods that turn imperfect scores into usable probabilities, and white-box probing methods that read internal representations when model access is available (Song et al., 2025; Liu et al., 2025; Cong et al., 2026).

In medical and EHR settings, this matters because LLMs are often overconfident when wrong, and their own stated confidence usually separates correct from incorrect answers only weakly (Omar et al., 2025; Omar et al., 2024; Savage et al., 2024; Naderi et al., 2025). EHR-specific studies add an important nuance: when the task is evidence retrieval or summarization from patient records, confidence can correlate meaningfully with faithfulness, but only when the score is tied to record-grounded evidence or stepwise reasoning rather than generic fluency (Ahsan et al., 2023; Song et al., 2025).

Are explainable algorithms available to estimate confidence for LLM outputs in EHR-based question answering?

Requires at least 5 papers that directly answer your question. Try adjusting your query to find more papers.

FIGURE 1 Research consensus on explainable confidence methods for EHR QA

The consensus is yes, but the strongest current methods are hybrid systems that combine interpretable evidence checks with calibration, not a single universal confidence score.

2. Methods

This review synthesizes the provided Deep Search corpus on LLM confidence estimation, calibration, explainability, hallucination detection, retrieval grounding, and EHR-focused medical question answering. The search strategy spanned general uncertainty-quantification papers, medical QA calibration studies, EHR-grounded QA systems, selective-answering frameworks, and explainability methods relevant to clinical use.

Search Strategy




FIGURE 2 Search funnel for included confidence-estimation literature

The strategy combined broad LLM uncertainty work with targeted EHR, medical QA, and evidence-grounding papers to isolate explainable confidence algorithms.

3. Results

3.1 Key Papers

A small set of papers anchors this literature because they define the main algorithmic options and expose their clinical limitations. Semantic entropy established meaning-level uncertainty detection for generative QA, Savage et al. compared uncertainty proxies directly on medical diagnosis and treatment tasks, and TrustEHRAgent translated confidence estimation into an EHR-specific agent with reliability-aware evaluation (Farquhar et al., 2024; Savage et al., 2024; Song et al., 2025).

Paper	Summary
(Farquhar et al., 2024)	Meaning-level entropy detects confabulations (Farquhar et al., 2024)
(Savage et al., 2024)	Sample consistency best in medical QA (Savage et al., 2024)
(Song et al., 2025)	Stepwise confidence for EHR agents (Song et al., 2025)

FIGURE 3 Foundational papers and their main contributions

3.2 Algorithm Families

The simplest explainable method is verbalized confidence: ask the LLM to output an answer and a probability or confidence score. This is black-box compatible and often easy to deploy, but in medical QA it is usually overconfident and poorly discriminative unless carefully prompted or fine-tuned (Yang et al., 2024; Tian et al., 2023; Savage et al., 2024; Tahermazandarani et al., 2026).

Consistency-based methods generate multiple answers and score confidence from agreement across them. Across general QA, medical QA, and biomedical NLP, these methods repeatedly outperform simple verbalized confidence on discrimination, especially when agreement is measured semantically rather than string-matched (Lyu et al., 2024; Savage et al., 2024; De Oliveira et al., 2025).

Semantic-uncertainty methods improve explainability by grouping sampled outputs by meaning rather than wording. Semantic entropy detects confabulations across tasks without task-specific tuning, and related medical adaptations such as word-sequence entropy improve free-form medical QA by weighting semantic relevance at both token and sequence levels (Farquhar et al., 2024; Wang et al., 2024).

Verification-based methods make the confidence score interpretable because they expose what failed. Two-phase Verification measures whether answers to generated verification questions stay consistent with the model's explanation, while stable-explanation methods estimate confidence from how stable the explanation distribution is across samples (Wu et al., 2024; Becker & Soatto, 2024).

3.3 EHR And Evidence Grounding

For EHR-based QA, evidence-grounded confidence is the most directly explainable family because the score can be attached to retrieved notes, contradictions, and missing evidence. In EHR evidence retrieval, model confidence correlated strongly with faithful summaries, suggesting that grounded record-level tasks are more amenable to useful confidence scoring than generic medical trivia (Ahsan et al., 2023).

Recent EHR systems operationalize this idea in different ways. TrustEHRAgent estimates confidence at each reasoning step and evaluates performance with HCAcc@k%, a metric that rewards answering accurately only above a reliability threshold (Song et al., 2025). MedConf constructs interpretable confidence from supporting, missing, and contradictory relations between available evidence and the diagnosis, and it stays robust when information is incomplete or irrelevant (Ren et al., 2026). EHR diagnostic prediction studies also show that hybrid feature-based calibrators using LLM embeddings plus XGBoost outperform both verbalized confidence and token logits on structured EHR tasks (Gao et al., 2024).

Atomic fact-checking is another explainable route. Medical RAG systems that decompose an answer into atomic facts and verify each claim against guidelines or EHR-derived propositions improve factuality while yielding fact-wise confidence and traceability (Vladika et al., 2025; Zeba et al., 2025).

3.4 Calibration, Guarantees, And White-Box Methods

Raw confidence scores are rarely good enough for deployment; post-hoc calibration is often necessary. Across biomedical NLP tasks, isotonic regression and histogram binning reduced calibration errors substantially, and multicalibration further improved both score error and calibration by enforcing reliability across intersecting prompt groups rather than only on average (De Oliveira et al., 2025; Detommaso et al., 2024).

If formal risk control is needed, conformal methods are the clearest option. Enhanced conformal prediction for medical MCQA uses self-consistency frequencies to form prediction sets with user-specified marginal error guarantees, while CIC converts arbitrary uncertainty scores into selective-answering thresholds with finite-sample control of accepted-answer error (Ke et al., 2025; Dong & Shinnou, 2026; Ye, 2025).

When model internals are accessible, probing methods often perform best. Hidden-state probes, activation-based classifiers, InternallInspector, and representation-stability methods consistently beat black-box heuristics in generic uncertainty benchmarks, and medical variants such as AU-Probe and Trust-Abstain show that a lightweight classifier over internal or trace-derived features can outperform self-reported confidence with low latency (Mahaut et al., 2024; Beigi et al., 2024; Khanmohammadi et al., 2025; Liu et al., 2026; D'Urso, 2026).

LLM confidence research has progressed from self-reported scores to grounded, calibrated, and selective systems.




FIGURE 4 Timeline of confidence-estimation research; larger markers indicate more citations

The timeline shows a clear shift from generic black-box elicitation toward medical-domain calibration, formal guarantees, and EHR-grounded verification pipelines.

Based on the included papers, a few authors and journals dominate visibility in this corpus.

Type	Name	Papers
Author	Thomas Savage	[4d9d7ac77e6c5973b214437b48a5b915][544f18a43c9c5958940fb0d35deb18a9] [66ddb2e8948d5fe2bcd09070181be91][4ca8ff75b97e57cfa11c2405b26dc5e9]
Author	Jimeng Sun	[be00016e4a2654879a53cfd508dfe9f9][4ca8ff75b97e57cfa11c2405b26dc5e9] [b11ec8e1b885703971ad7253e9a5496]
Author	Eyal Klang	[9323b11e06775363bf49e350f66c9eaa][22c058329d395aca8299c711b297154e]
Journal	ArXiv	[9656a4c7bf7d5eee974c64d83220a9ed][03b4a28a7ebd5949b584a172c9a8f34e] [2291ace7a4c85663a0c794a75e4de0f2][30ea5a0cb6b8541ba175a690c7ee2cba] [ad3e5341ac045c0094014384abcb5015]
Journal	JAMIA	[4d9d7ac77e6c5973b214437b48a5b915][e768d463c3905353a67fd6b017e17bf9]
Journal	NPJ Digital Medicine	[e55016ea68de540294713aeb397c8c4e][dd265d40419459ef9f57f3da69f458e9] [6637ea97176d5c61b32867954d6f50af]

FIGURE 5 Authors and journals appearing most often in corpus

4. Discussion

The strongest practical conclusion is that explainable confidence for EHR-based QA should be built around evidence and behavior, not only around the final answer token. Verbalized confidence is attractive because it is cheap and understandable, but multiple medical studies show that it remains systematically overconfident and varies substantially by specialty, task type, and prompting protocol (Boie et al., 2026; Savage et al., 2024; Naderi et al., 2025; Kim & Kang, 2026).

Consistency and semantic-dispersion methods have the best cross-paper support among black-box approaches, but they are computationally expensive because they require multiple generations (Savage et al., 2024; Lin et al., 2023). That tradeoff explains why newer work tries to distill or approximate these signals into single-pass models, such as APRICOT, uncertainty distillation, or reasoning-trace classifiers (Ulmer et al., 2024; Hager et al., 2025; Testoni & Calixto, 2025).

For clinical deployment, the most convincing systems are those that combine confidence with abstention, retrieval, or verification. Evidence-grounded RAG, atomic fact checking, conformal selective answering, and confidence-aware multi-agent routing all turn confidence into an auditable workflow decision rather than an isolated number (Xie et al., 2026; Vladika et al., 2025; Dong & Shinnou, 2026; Elshaer & Rashed, 2025; Nair-Kanneganti et al., 2025).

Claim	Evidence Strength	Reasoning	Papers
Verbalized confidence alone is unreliable for medical QA	 Strong	Repeated medical benchmarks show overconfidence and weak separation of right vs wrong answers.	(Omar et al., 2025; Omar et al., 2024; Savage et al., 2024; Naderi et al., 2025)
Sample consistency / semantic entropy are the best black-box baselines	 Strong	Multiple studies show better discrimination than verbalized or token-only scores.	(Savage et al., 2024; Farquhar et al., 2024; Lyu et al., 2024)
Evidence-grounded and verification-based scores are most explainable for EHR QA	 Moderate	They tie confidence to support, contradiction, or reasoning consistency that clinicians can inspect.	(Ahsan et al., 2023; Vladika et al., 2025; Ren et al., 2026; Wu et al., 2024)
White-box probing often outperforms black-box heuristics	 Moderate	Strong generic evidence, but less direct EHR-specific validation.	(Mahaut et al., 2024; Khanmohammadi et al., 2025; Beigi et al., 2024; Liu et al., 2026)

Claim	Evidence Strength	Reasoning	Papers
Formal risk guarantees for open-ended EHR QA remain limited	 Weak	Conformal and selective methods are promising, but mostly validated on MCQA or generic QA.	(Ke et al., 2025; Dong & Shinnou, 2026; Ye, 2025)

FIGURE 6 Key claims and evidence across confidence-estimation methods

5. Conclusion

The best answer to the question is that explainable confidence scores for LLM outputs in EHR-based question answering do exist, but the strongest options are hybrid algorithms. They combine interpretable evidence support, contradiction checks, or reasoning verification with calibration or abstention, rather than trusting the model’s own confidence statement alone (Ren et al., 2026; Song et al., 2025; Vladika et al., 2025).

Research Gaps

The literature now covers several algorithm families, but coverage is uneven across EHR tasks, confidence mechanisms, and deployment guarantees. Most work still evaluates multiple-choice medical QA or generic factual QA rather than longitudinal, messy, incomplete EHR dialogue.

	EHR tasks	Open-ended QA	Selective abstain	Clinician-facing
Verbalized	GAP	GAP	GAP	GAP
Consistency	GAP	GAP	GAP	GAP
Evidence-grounded	GAP	GAP	GAP	GAP
Internal probes	GAP	GAP	GAP	GAP
Conformal	GAP	GAP	GAP	GAP

FIGURE 7 Coverage gaps across confidence methods and clinical settings

The biggest gap is formal, clinician-auditable confidence for open-ended EHR reasoning. Conformal methods offer guarantees, but current medical applications focus mainly on multiple-choice prediction sets rather than long-form patient-record answers (Ke et al., 2025; Ye, 2025). Conversely, EHR-grounded systems are interpretable and practically useful, but they usually report empirical gains without the same level of statistical risk control (Song et al., 2025; Ren et al., 2026; Ibrahim et al., 2025). Another persistent gap is user-facing explanation design: clinicians need confidence tied to specific symptoms, labs, missing evidence, and contradictory notes, but only a subset of systems currently expose that structure explicitly (Nourmohammadi et al., 2026; Atf et al., 2025).

Open Research Questions

The next phase of the field is moving from “can we detect uncertainty?” to “can we express uncertainty in ways clinicians can safely act on?”

Question	Why
How can conformal risk control be extended from medical MCQA to open-ended EHR-grounded generation?	Current guarantees largely stop at prediction sets, while real EHR assistants generate free text and cited rationales.
Which single-pass interpretable features best approximate expensive sampling-based confidence for EHR QA?	Sampling works well but is slow; scalable hospital deployment needs cheaper surrogates with similar reliability.
How should confidence explanations link supporting, missing, and contradictory EHR evidence for clinicians?	Usable confidence in practice likely depends on evidence structure, not just a scalar probability.

FIGURE 8 Open research questions for explainable EHR confidence

The most important shift is from opaque self-confidence to evidence-linked confidence; the biggest open question is how to give that confidence formal guarantees in open-ended EHR QA.

These search results were found and analyzed using Consensus, an AI-powered search engine for research. Try it at <https://consensus.app>. © 2026 Consensus NLP, Inc. Personal, non-commercial use only; redistribution requires copyright holders' consent.

References

- Ahsan, H., McInerney, D. J., Kim, J., Potter, C., Young, G., Amir, S., & Wallace, B. C. (2023). Retrieving Evidence from EHRs with LLMs: Possibilities and Challenges. *Proceedings of machine learning research*, 248, 489 - 505. <https://doi.org/10.48550/arxiv.2309.04550>
- Atf, Z., Safavi-Naini, S. A. A., Lewis, P. R., Mahjoubfar, A., Naderi, N., Savage, T., & Soroush, A. (2025). The challenge of uncertainty quantification of large language models in medicine. *ArXiv, abs/2504.05278*. <https://doi.org/10.48550/arxiv.2504.05278>
- Becker, E., & Soatto, S. (2024). Cycles of Thought: Measuring LLM Confidence through Stable Explanations. *ArXiv, abs/2406.03441*. <https://doi.org/10.48550/arxiv.2406.03441>
- Beigi, M., Shen, Y., Yang, R., Lin, Z., Wang, Q., Mohan, A., He, J., Jin, M., Lu, C.-T., & Huang, L. (2024). InternallInspector I2: Robust Confidence Estimation in LLMs through Internal States. *ArXiv, abs/2406.12053*. <https://doi.org/10.48550/arxiv.2406.12053>
- Boie, S., Reis, F., Frey, N., Grünwald, E., & Balzer, F. (2026). Calibration of Self-Reported Confidence and Accuracy of Large Language Models in Medical Question Answering. *Journal of Medical Systems*, 50. <https://doi.org/10.1007/s10916-026-02430-0>
- Cong, L., Hahn, S., Gombert, S., Camus, L., Drachsler, H., & Kroehne, U. (2026). Confidence Estimation in Automatic Short Answer Grading with LLMs. 110-123. <https://doi.org/10.48550/arxiv.2605.00200>

- D'Urso, E. (2026). Hidden-Uncertainty Assessment via Non-Verbalized Signatures for LLM Medical QA. *IEEE Access*, 14, 58744-58759. <https://doi.org/10.1109/access.2026.3682338>
- De Oliveira, R. D., Garber, M., Gwinnutt, J., Rashidi, E., Hwang, J.-H., Gilmour, W., Nanavati, J., Abidine, K. Z. E., & Mack, C. (2025). A study of calibration as a measurement of trustworthiness of large language models in biomedical natural language processing. *JAMIA Open*, 8. <https://doi.org/10.1093/jamiaopen/ooaf058>
- Detommaso, G., Bertran Lopez, M., Fogliato, R., & Roth, A. (2024). Multicalibration for Confidence Scoring in LLMs. *ArXiv, abs/2404.04689*. <https://doi.org/10.48550/arxiv.2404.04689>
- Dong, S., & Shinnou, H. (2026). Uncertainty-Aware Abstention in Large Language Models with Provable Alignment Guarantees. *ArXiv, abs/2607.04430*. <https://doi.org/10.48550/arxiv.2607.04430>
- Elshaer, Z., & Rashed, E. (2025). CURE: Confidence-driven Unified Reasoning Ensemble Framework for Medical Question Answering. *Big Data Cogn. Comput.*, 9, 299. <https://doi.org/10.48550/arxiv.2510.14353>
- Farquhar, S., Kossen, J., Kuhn, L., & Gal, Y. (2024). Detecting hallucinations in large language models using semantic entropy. *Nature*, 630, 625 - 630. <https://doi.org/10.1038/s41586-024-07421-0>
- Gao, Y., Myers, S., Chen, S., Dligach, D., Miller, T. A., Bitterman, D., Chen, G., Mayampurath, A., Churpek, M., & Afshar, M. (2024). Uncertainty estimation in diagnosis generation from large language models: next-word probability is not pre-test probability. *JAMIA Open*, 8. <https://doi.org/10.1093/jamiaopen/ooae154>
- Hager, S., Mueller, D., Duh, K., & Andrews, N. (2025). Uncertainty Distillation: Teaching Language Models to Express Semantic Confidence. *ArXiv, abs/2503.14749*. <https://doi.org/10.48550/arxiv.2503.14749>
- Ibrahim, A., Khalili, A., Arabi, M., Sattar, A., Hosseini, A., & Serag, A. (2025). MERA: Medical Electronic Records Assistant. *Mach. Learn. Knowl. Extr.*, 7, 73. <https://doi.org/10.3390/make7030073>
- Ke, Y., Lin, H., Ruan, Y., Tang, J., & Li, L. (2025). Correctness Coverage Evaluation for Medical Multiple-Choice Question Answering Based on the Enhanced Conformal Prediction Framework. *Mathematics*. <https://doi.org/10.3390/math13091538>
- Khanmohammadi, R., Miah, E., Mardikoraem, M., Kaur, S., Brugere, I., Smiley, C. H., Thind, K., & Ghassemi, M. M. (2025). Calibrating LLM Confidence by Probing Perturbed Representation Stability. 10448-10514. <https://doi.org/10.48550/arxiv.2505.21772>
- Kim, H., & Kang, P. (2026). Same Answer, Different Confidence: Protocol Sensitivity in LLM Confidence Calibration. <https://doi.org/10.48550/arxiv.2605.27752>
- Lin, Z., Trivedi, S., & Sun, J. (2023). Generating with Confidence: Uncertainty Quantification for Black-box Large Language Models. *Trans. Mach. Learn. Res.*, 2024. <https://doi.org/10.48550/arxiv.2305.19187>
- Liu, X., Chen, T., Da, L., Chen, C., Lin, Z.-Y., & Wei, H. (2025). Uncertainty Quantification and Confidence Calibration in Large Language Models: A Survey. *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*. <https://doi.org/10.1145/3711896.3736569>
- Liu, Y., Liu, Y., Mbuvi, P., Li, Z., Yao, R., Lim, G., & Wang, D. (2026). Mind the Ambiguity: Aleatoric Uncertainty Quantification in LLMs for Safe Medical Question Answering. *Proceedings of the ACM Web Conference 2026*. <https://doi.org/10.1145/3774904.3793027>
- Lyu, Q., Shridhar, K., Malaviya, C., Zhang, L., Elazar, Y., Tandon, N., Apidianaki, M., Sachan, M., & Callison-Burch, C. (2024). Calibrating Large Language Models with Sample Consistency. 19260-19268. <https://doi.org/10.48550/arxiv.2402.13904>

- Mahaut, M., Aina, L., Czarnowska, P., Hardalov, M., Müller, T., & Marquez, L. (2024). Factual Confidence of LLMs: on Reliability and Robustness of Current Estimators. *ArXiv, abs/2406.13415*. <https://doi.org/10.18653/v1/2024.acl-long.250>
- Naderi, N., Safavi-Naini, S. A. A., Savage, T., Atf, Z., Lewis, P. R., Nadkarni, G., & Soroush, A. (2025). Across generations, sizes, and types, large language models poorly report self-confidence in gastroenterology clinical reasoning tasks. *npj Gut and Liver, 3*. <https://doi.org/10.1038/s44355-026-00053-3>
- Nair-Kanneganti, A., Chan, T. J., Goldfinger, S., Mackay, E., Anthony, B., & Pouch, A. M. (2025). Increasing LLM response trustworthiness using voting ensembles. *ArXiv, abs/2510.04048*. <https://doi.org/10.48550/arxiv.2510.04048>
- Nourmohammadi, N., Hossain, M. M., Han, T., Ara, S., & Shamszaman, Z. U. (2026). A Multi-Agent Framework for Medical AI: Leveraging Fine-Tuned GPT, LLaMA, and DeepSeek R1 for Evidence-Based and Bias-Aware Clinical Query Processing. *ArXiv, abs/2602.14158*. <https://doi.org/10.48550/arxiv.2602.14158>
- Omar, M., Agbareia, R., Glicksberg, B. S., Nadkarni, G., & Klang, E. (2025). Benchmarking the Confidence of Large Language Models in Answering Clinical Questions: Cross-Sectional Evaluation Study. *JMIR Medical Informatics, 13*. <https://doi.org/10.2196/66917>
- Omar, M., Glicksberg, B. S., Nadkarni, G., & Klang, E. (2024). Overconfident AI? Benchmarking LLM Self-Assessment in Clinical Scenarios. <https://doi.org/10.1101/2024.08.11.24311810>
- Ren, Z., Zhan, Y., Liang, S., G., Yu, B., & Tao, D. (2026). Towards Reliable Medical LLMs: Benchmarking and Enhancing Confidence Estimation of Large Language Models in Medical Consultation. *ArXiv, abs/2601.15645*. <https://doi.org/10.48550/arxiv.2601.15645>
- Savage, T., Wang, J., Gallo, R., Boukil, A., Patel, V., Safavi-Naini, S. A. A., Soroush, A., & Chen, J. H. (2024). Large language model uncertainty proxies: discrimination and calibration for medical diagnosis and treatment. *Journal of the American Medical Informatics Association : JAMIA*. <https://doi.org/10.1093/jamia/ocae254>
- Savage, T. R., Wang, J., Gallo, R., Patel, V., Soroush, A., & Chen, J. H. (2024). Large Language Model Uncertainty Measurement and Calibration for Medical Diagnosis and Treatment. <https://doi.org/10.1101/2024.06.06.24308399>
- Song, Y., Jeong, M., & Sung, M. (2025). Trustworthy Agents for Electronic Health Records through Confidence Estimation. *ArXiv, abs/2508.19096*. <https://doi.org/10.48550/arxiv.2508.19096>
- Tahermazandarani, M., Mahmood, A., Islam, F., & Sheng, Q. Z. (2026). When Confidence Fails: Overconfidence in LLMS Under Uncertainty and Missing Clinical Information. *2026 IEEE International Conference on Digital Health (ICDH)*, 315-323. <https://doi.org/10.1109/icdh72779.2026.00043>
- Testoni, A., & Calixto, I. (2025). Different Questions, Different Models: Fine-Grained Evaluation of Uncertainty and Calibration in Clinical QA with LLMs. *ArXiv, abs/2506.10769*. <https://doi.org/10.48550/arxiv.2506.10769>
- Tian, K., Mitchell, E., Zhou, A., Sharma, A., Rafailov, R., Yao, H., Finn, C., & Manning, C. D. (2023). Just Ask for Calibration: Strategies for Eliciting Calibrated Confidence Scores from Language Models Fine-Tuned with Human Feedback. *ArXiv, abs/2305.14975*. <https://doi.org/10.48550/arxiv.2305.14975>
- Ulmer, D., Gubri, M., Lee, H., Yun, S., & Oh, S. J. (2024). Calibrating Large Language Models Using Their Generations Only. 15440-15459. <https://doi.org/10.48550/arxiv.2403.05973>

Vladika, J., Domres, A., Nguyen, M., Moser, R., Nano, J., Busch, F., Adams, L., Bressemer, K., Bernhardt, D., Combs, S. E., Borm, K., Matthes, F., & Peeken, J. (2025). Improving Reliability and Explainability of Medical Question Answering through Atomic Fact Checking in Retrieval-Augmented LLMs. *ArXiv, abs/2505.24830*.

<https://doi.org/10.48550/arxiv.2505.24830>

Wang, Z., Duan, J., Yuan, C., Chen, Q., Chen, T., Yao, H., Zhang, Y., Wang, R., Xu, K., & Shi, X. (2024). Word-Sequence Entropy: Towards Uncertainty Estimation in Free-Form Medical Question Answering Applications and Beyond. *Eng. Appl. Artif. Intell.*, 139, 109553. <https://doi.org/10.48550/arxiv.2402.14259>

Wu, J., Yu, Y., & Zhou, H.-Y. (2024). Uncertainty Estimation of Large Language Models in Medical Question Answering. *ArXiv, abs/2407.08662*. <https://doi.org/10.48550/arxiv.2407.08662>

Xie, X., Liu, X., & Yao, B. Z. (2026). Know When You're Wrong: Aligning Confidence with Correctness for LLM Error Detection. *ArXiv, abs/2603.06604*. <https://doi.org/10.48550/arxiv.2603.06604>

Yang, D., Tsai, Y.-H., & Yamada, M. (2024). On Verbalized Confidence Scores for LLMs. *ArXiv, abs/2412.14737*. <https://doi.org/10.48550/arxiv.2412.14737>

Ye, Y.-C. (2025). Conformal P-Value in Multiple-Choice Question Answering Tasks with Provable Risk Control. *ArXiv, abs/2508.10022*. <https://doi.org/10.48550/arxiv.2508.10022>

Zeba, M., Mamun, A. A., Tithee, K. J., Sutradhar, D., Raiaan, M. A. K., Mukta, S., Mohamed, R. E., Islam, M. R., Sebastian, Y., Hussain, M., & Azam, S. (2025). Mitigating Hallucinations in Healthcare LLMs with Granular Fact-Checking and Domain-Specific Adaptation. *Expert Syst. Appl.*, 329, 132966.

<https://doi.org/10.1016/j.eswa.2026.132966>